

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import os
import scipy as sp
import warnings
warnings.filterwarnings("ignore")
%matplotlib inline
```

```
data = pd.read_csv('diabetes.csv')
```

```
data.head()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
data.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFu	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471011	33.234454	0.348160
std	3.369578	41.985160	19.355807	15.952218	115.244002	7.884160	0.471011	11.969603	0.477297
min	0.000000	64.000000	0.000000	0.000000	0.000000	0.000000	0.000000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.167350	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.369125	33.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.672277	36.000000	0.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.425000	81.000000	1.000000

Saved successfully!

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                    768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                    768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
data.shape
```

(768, 9)

```
data.value_counts()
```

Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome	
0	57	60	0	0	21.7	0.735	67	0	1
	67	76	0	0	45.3	0.194	46	0	1
5	103	108	37	0	39.2	0.305	65	0	1
	104	74	0	0	28.8	0.153	48	0	1
	105	72	29	325	36.9	0.159	28	0	1

```
2      84      50      23      76      30.4 0.968      21 0      ..
      85      65      0      0      39.6 0.930      27 0      1
      87      0      23      0      28.9 0.773      25 0      1
      58      16      52      32.7 0.166      25 0      1
17     163      72      41      114     40.9 0.817      47 1      1
Length: 768, dtype: int64
```

data.dtypes

```
Pregnancies      int64
Glucose           int64
BloodPressure     int64
SkinThickness     int64
Insulin           int64
BMI               float64
DiabetesPedigreeFunction float64
Age              int64
Outcome           int64
dtype: object
```

data.columns

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
       'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

data.isnull().sum()

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction 0
Age              0
```

Saved successfully! X

data.isnull().any()

```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction False
Age              False
Outcome           False
dtype: bool
```

data.isnull().all()

```
Pregnancies      False
Glucose           False
BloodPressure     False
SkinThickness     False
Insulin           False
BMI               False
DiabetesPedigreeFunction False
Age              False
Outcome           False
dtype: bool
```


data.corr()

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabete
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	


```
plt.figure(figsize = (12,10))
```

```
sns.heatmap(data.corr(), annot =True)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fba5d871c70>

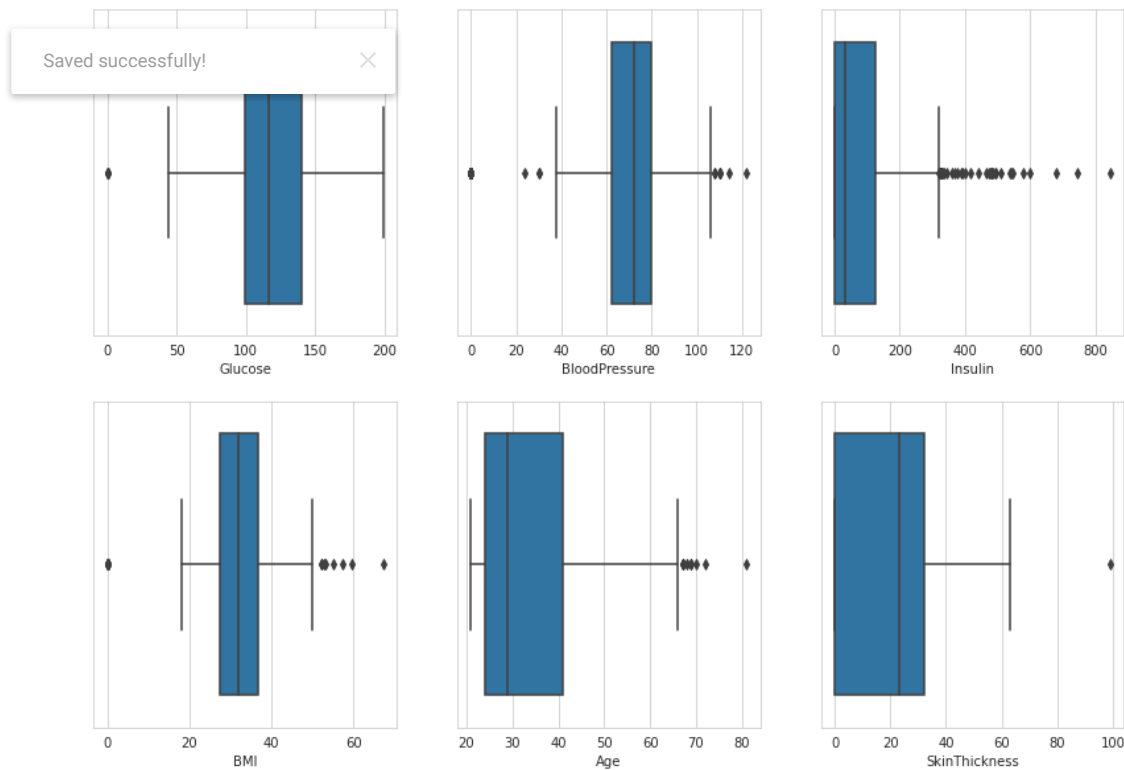


```
data.hist(figsize=(18,12))  
plt.show()
```



```
plt.figure(figsize=(14,10))
sns.set_style(style='whitegrid')
plt.subplot(2,3,1)
sns.boxplot(x='Glucose',data=data)
plt.subplot(2,3,2)
sns.boxplot(x='BloodPressure',data=data)
plt.subplot(2,3,3)
sns.boxplot(x='Insulin',data=data)
plt.subplot(2,3,4)
sns.boxplot(x='BMI',data=data)
plt.subplot(2,3,5)
sns.boxplot(x='Age',data=data)
plt.subplot(2,3,6)
sns.boxplot(x='SkinThickness',data=data)
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7fba5d2746a0>




```
mean_col = ['Glucose','BloodPressure','Insulin','Age','Outcome','BMI']
sns.pairplot(data[mean_col],palette='Accent')
```

```
<seaborn.axisgrid.PairGrid at 0x7fba5d1a57c0>
```




```
sns.boxplot(x='Outcome',y='Insulin',data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fba5c8976d0>
```




```
sns.regplot(x='BMI', y= 'Glucose', data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fba5c7e2df0>
```



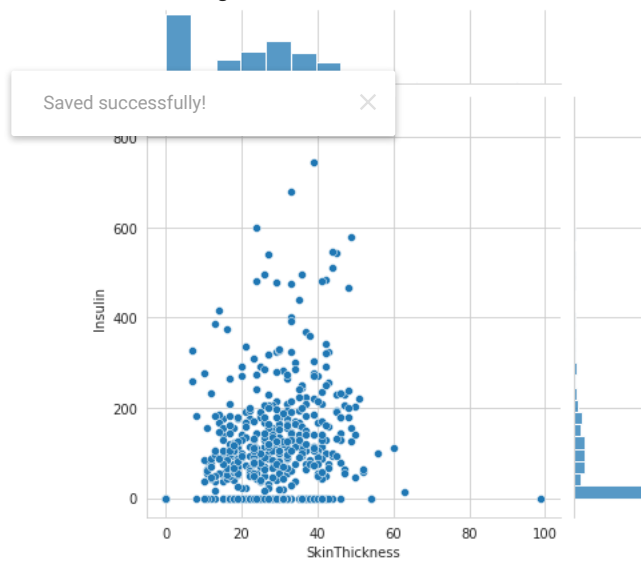
```
sns.scatterplot(x='Glucose', y= 'Insulin', data=data)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fba5c7cf040>
```



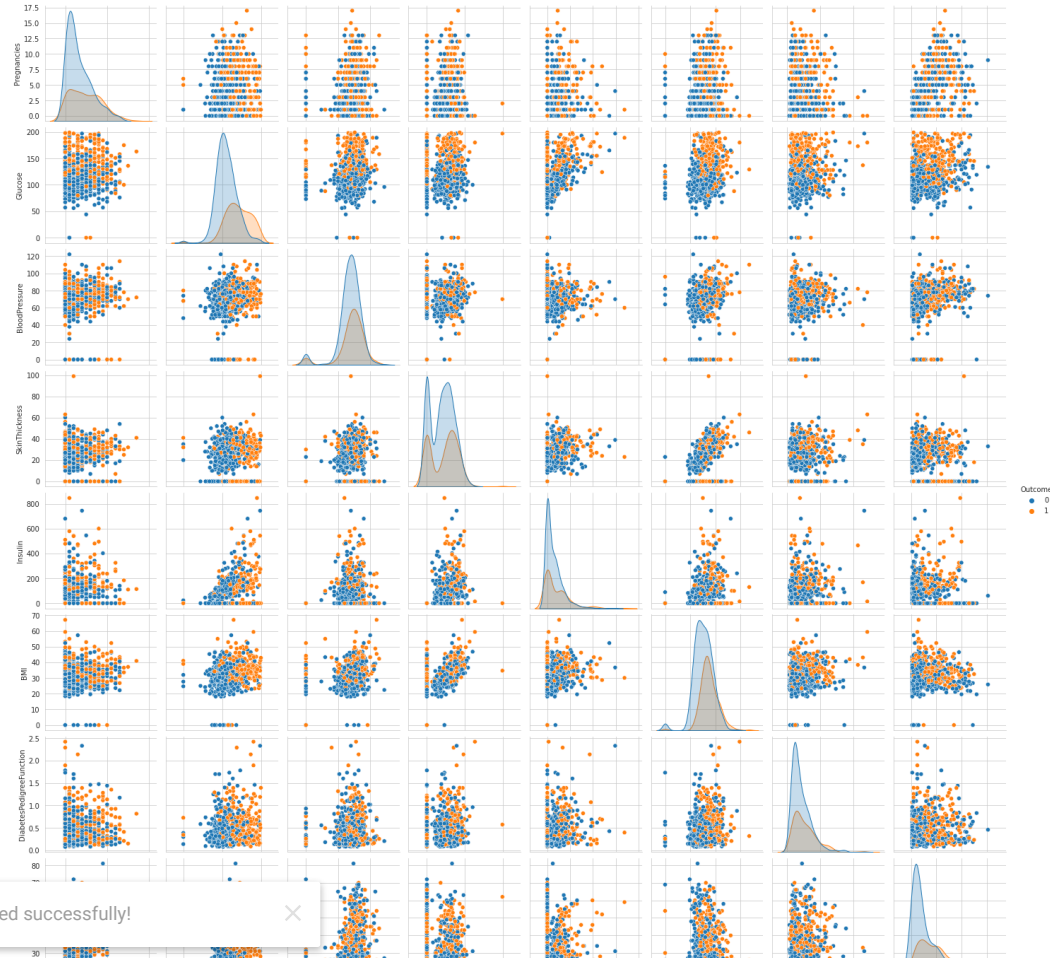
```
sns.jointplot(x='SkinThickness', y= 'Insulin', data=data)
```

```
<seaborn.axisgrid.JointGrid at 0x7fba5c724700>
```



```
sns.pairplot(data,hue='Outcome')
```


&lt;seaborn.axisgrid.PairGrid at 0x7fba5c6d1760&gt;



Saved successfully!

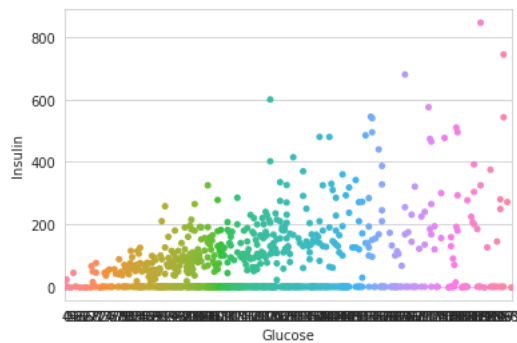
```
sns.lineplot(x='Glucose', y= 'Insulin', data=data)
```

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fba591a5550&gt;

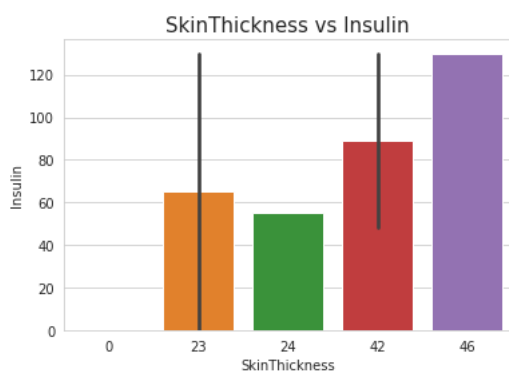


```
sns.swarmplot(x='Glucose', y= 'Insulin', data=data)
```


&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7fba590c9490&gt;



```
sns.barplot(x="SkinThickness", y="Insulin", data=data[170:180])
plt.title("SkinThickness vs Insulin",fontsize=15)
plt.xlabel("SkinThickness")
plt.ylabel("Insulin")
plt.show()
plt.style.use("ggplot")
```



```
plt.style.use("default")
plt.figure(figsize=(5,5))
sns.barplot(x="Glucose", y="Insulin", data=data[170:180])
plt.title("Glucose vs Insulin",fontsize=15)
plt.xlabel("Glucose")
plt.ylabel("Insulin")
plt.show()
```



```
x = data.drop(columns = 'Outcome')

# Getting Predicting Value
y = data['Outcome']

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.2,random_state=0)

print(len(x_train))
print(len(x_test))
print(len(y_train))
print(len(y_test))
```

```
614
154
614
154
```



```
from sklearn.linear_model import LogisticRegression
reg = LogisticRegression()
reg.fit(x_train,y_train)

LogisticRegression()

y_pred=reg.predict(x_test)
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
print("Classification Report is:\n",classification_report(y_test,y_pred))
print("Confusion Matrix:\n",confusion_matrix(y_test,y_pred))
print("Training Score:\n",reg.score(x_train,y_train)*100)
print("Mean Squared Error:\n",mean_squared_error(y_test,y_pred))
print("R2 score is:\n",r2_score(y_test,y_pred))
```

Classification Report is:

	precision	recall	f1-score	support
0	0.84	0.92	0.88	107
1	0.76	0.62	0.68	47
accuracy			0.82	154
macro avg	0.80	0.77	0.78	154
weighted avg	0.82	0.82	0.82	154

Confusion Matrix:  
[[98 9]  
[18 29]]  
Training Score:  
77.19869706840392  
Mean Squared Error:  
0.17532467532467533  
R2 score is:  
0.1731954662954862

Saved successfully!

82.46753246753246

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=7)

knn.fit(x_train,y_train)

KNeighborsClassifier(n_neighbors=7)

y_pred=knn.predict(x_test)
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
print("Classification Report is:\n",classification_report(y_test,y_pred))
print("Confusion Matrix:\n",confusion_matrix(y_test,y_pred))
print("Training Score:\n",knn.score(x_train,y_train)*100)
print("Mean Squared Error:\n",mean_squared_error(y_test,y_pred))
print("R2 score is:\n",r2_score(y_test,y_pred))
```

Classification Report is:

	precision	recall	f1-score	support
0	0.82	0.84	0.83	107
1	0.61	0.57	0.59	47
accuracy			0.76	154
macro avg	0.72	0.71	0.71	154
weighted avg	0.76	0.76	0.76	154

Confusion Matrix:  
[[90 17]  
[20 27]]  
Training Score:  
78.17589576547232  
Mean Squared Error:  
0.24025974025974026  
R2 score is:  
-0.13302843507655582

```
print(accuracy_score(y_test,y_pred)*100)
```

```
75.97402597402598
```

```
from sklearn.svm import SVC
```

```
svc = SVC()
```

```
svc.fit(x_train, y_train)
```

```
SVC()
```

```
y_pred=svc.predict(x_test)
```

```
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
```

```
from sklearn.metrics import r2_score
```

```
from sklearn.metrics import mean_squared_error
```

```
print("Classification Report is:\n",classification_report(y_test,y_pred))
```

```
print("Confusion Matrix:\n",confusion_matrix(y_test,y_pred))
```

```
print("Training Score:\n",svc.score(x_train,y_train)*100)
```

```
print("Mean Squared Error:\n",mean_squared_error(y_test,y_pred))
```

```
print("R2 score is:\n",r2_score(y_test,y_pred))
```

```
Classification Report is:
```

```
precision    recall  f1-score   support
```

```
0           0.81     0.92     0.86     107
```

```
1           0.73     0.51     0.60     47
```

```
accuracy          0.79     154
```

```
macro avg         0.77     0.71     0.73     154
```

```
weighted avg      0.78     0.79     0.78     154
```

```
Confusion Matrix:
```

```
[[98  9]
```

```
 [ 9 47]]
```

Saved successfully!



```
Mean Squared Error:
```

```
0.2077922077922078
```

```
R2 score is:
```

```
0.020083515609465197
```

```
print(accuracy_score(y_test,y_pred)*100)
```

```
79.22077922077922
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
gbc=GradientBoostingClassifier()
```

```
gbc.fit(x_train,y_train)
```

```
GradientBoostingClassifier()
```

```
y_pred=gbc.predict(x_test)
```

```
from sklearn.metrics import accuracy_score,classification_report,confusion_matrix
```

```
from sklearn.metrics import r2_score
```

```
from sklearn.metrics import mean_squared_error
```

```
print("Classification Report is:\n",classification_report(y_test,y_pred))
```

```
print("Confusion Matrix:\n",confusion_matrix(y_test,y_pred))
```

```
print("Training Score:\n",gbc.score(x_train,y_train)*100)
```

```
print("Mean Squared Error:\n",mean_squared_error(y_test,y_pred))
```

```
print("R2 score is:\n",r2_score(y_test,y_pred))
```

```
Classification Report is:
```

```
precision    recall  f1-score   support
```

```
0           0.87     0.86     0.86     107
```

```
1           0.69     0.70     0.69     47
```

```
accuracy          0.81     154
```

```
macro avg         0.78     0.78     0.78     154
```

```
weighted avg      0.81     0.81     0.81     154
```

```
Confusion Matrix:
```

```
[[92 15]
```

```
 [14 33]]
```

```
Training Score:  
91.85667752442997  
Mean Squared Error:  
0.18831168831168832  
R2 score is:  
0.11195068602107783
```

```
print(accuracy_score(y_test,y_pred)*100)
```

```
81.16883116883116
```

✓ 0s completed at 2:25 PM



Saved successfully!

